TECHNICAL SPECIFICATION

IFRS TAXONOMY & CURRENCY-AWARE NORMALIZATION

Status: Approved 2026-07-11 (design §9) — not yet implemented **Companions:** `docs/ifrs-currency-support-design.md` **Parent:** `docs/architecture.md` , `docs/tech_spec.md` **Audience:** Implementing engineers / agents

---

1. Scope

Three work packages, independently shippable in order:

- **WP1 — IFRS concept coverage** (closes the INFY blank screen; additive).
- **WP2 — debt semantics policy** (IFRS 16; **policy approved** — leases are not debt).
- **WP3 — currency-aware normalization** (WP3a guards + USD-preference; WP3b FX conversion). **Both in scope this cycle** (WP3b moved in 2026-07-11, design §9).

Plus cross-cutting **guards** (WP-G) that ship with WP1.

Nothing here changes the read-time derivation model, the period reconciler, or the IPC surface. The normalization engine is already taxonomy-agnostic (`orchestrator.rs` → `concept_map::metric_for(&f.taxonomy, &f.concept)`), so WP1 is pure data addition to one file.

---

2. WP1 — IFRS concept coverage

2.1 Concept-map additions — `src-tauri/src/normalize/concept_map.rs`

Extend each `concepts_for(metric)` arm by appending `ifrs-full` candidates **after** the existing `us-gaap` (and `dei`) candidates. Ordering matters: US-GAAP first preserves current behavior for US filers; IFRS candidates are fallbacks that only fire for IFRS filers. Within IFRS, list the modern (IFRS 15/16-era) concept before the legacy one.

`Metric` arm	Append these `("ifrs-full", …)` candidates (in order)
`Revenue`	`RevenueFromContractsWithCustomers` , `Revenue`

`Metric` arm	Append these `("ifrs-full", …)` candidates (in order)
`CostOfRevenue`	`CostOfSales`
`GrossProfit`	`GrossProfit`
`OperatingIncome`	`ProfitLossFromOperatingActivities`
`NetIncome`	`ProfitLoss`
`EpsBasic`	`BasicEarningsLossPerShare`
`EpsDiluted`	`DilutedEarningsLossPerShare`
`SharesOutstandingBasic`	`WeightedAverageShares`
`SharesOutstandingDiluted`	`AdjustedWeightedAverageShares`
`CashAndEquivalents`	`CashAndCashEquivalents`
`TotalAssets`	`Assets`
`TotalLiabilities`	`Liabilities`
`TotalEquity`	`EquityAttributableToOwnersOfParent`, `Equity`
`CashFromOperations`	`CashFlowsFromUsedInOperatingActivities`
`CapitalExpenditures`	`PurchaseOfPropertyPlantAndEquipmentClassifiedAsInvestingActivities`
`DepreciationAmortization`	`DepreciationAndAmortisationExpense`
`PropertyPlantAndEquipmentNet`	`PropertyPlantAndEquipment`
`LongTermDebt` / `CurrentDebt`	**See WP2 — borrowings only (policy approved); mappings listed there**

Notes:

- `TotalEquity`: parent-attributable first, incl-NCI (`Equity`) fallback — mirrors the existing US-GAAP ordering (`StockholdersEquity` before the incl-NCI variant).
- `CapitalExpenditures`: `apply_sign` already `abs()`-normalizes this metric (`orchestrator.rs:430`), so the IFRS investing-outflow sign needs no special handling.
- Bank metrics (`NetInterestIncome`, `NoninterestIncome`, `InterestIncomeOperating`, `InterestExpense`) stay US-GAAP-only; annotate them explicitly for the coverage test (§5.1).
- No change to `metric_for` — it already iterates every candidate across all taxonomies.

2.2 Derivations — no change required

`free_cash_flow` (= `NetIncome` + `DepreciationAmortization` − `CapitalExpenditures`), the capex PP&E roll-forward fallback (`PropertyPlantAndEquipmentNet` + `DepreciationAmortization`), `operating_margin`, and the TTM/yield chain all resolve once the §2.1 metrics are populated. Verified against INFY's reported facts: every input concept is present.

2.3 Acceptance

INFY (CIK 0001067491) renders non-empty Income / Balance / Cashflow / Valuation series. Re-ingest, then confirm `normalized_fact` count for the CIK jumps from 9 to the expected multi-metric set.

---

3. WP2 — debt semantics (IFRS 16)

Policy approved 2026-07-11 (design §5.3, §9): **economic equivalence** — map interest-bearing borrowings only; exclude lease liabilities, matching the US-GAAP treatment.

Append to the debt arms:

`Metric` arm	**Append `("ifrs-full", …)`**
`LongTermDebt`	`NoncurrentBorrowings`, `Borrowings`
`CurrentDebt`	`CurrentBorrowings`

Do **not** map `LeaseLiabilities` / `NoncurrentLeaseLiabilities` / `CurrentLeaseLiabilities` to any debt metric. For Infosys this correctly yields ~zero debt (`Borrowings` : 1 fact; no material borrowings). A distinct `LeaseLiabilities` metric is deferred (design §5.2 Option C).

---

4. WP3 — currency-aware normalization

4.1 WP3a — guards now (cheap, honest)

Split currency detection from scaling — `src-tauri/src/sources/companyfacts.rs` :

```
/// Monetary units are scaled ×1e6; `shares` ×1; the currency (if any) is
/// returned separately so downstream can resolve it to USD.
fn classify_unit(unit: &str) -> UnitClass {
    match unit {
        "shares"           => UnitClass::Count,
        "pure"             => UnitClass::Ratio,          // ×1e6, no currency
        u if u.contains('/') => UnitClass::PerShare(currency_of(u)), // e.g. USD/shares, IN
        u                  => UnitClass::Monetary(currency_of(u)),   // USD, INR, EUR, …
    }
}
```

`scale_to_micro` keeps its ×1e6 / ×1 behavior; it stops being the place that implicitly declares "everything is USD."

USD-preference selection + reject/diagnostic — `src-tauri/src/pipeline/orchestrator.rs`:

In `best_annual` / `best_instant` selection and in the quarterly path, when the winning fact's currency ≠ USD:

1. If a USD-unit sibling fact exists for the same (metric, period), prefer it (Option A).
2. Else skip the fact and emit a `Severity::Warn` ingestion event (Option C): *"Skipped {metric} for FY{y}: reported in {ccy}, USD conversion not yet available."*

Stop hardcoding `unit: "USD"` in the quarterly branch (`orchestrator.rs:185`); flow the resolved currency (which, after this ladder, is always USD or the fact is skipped).

4.2 WP3b — FX conversion (unlocks general FPI support)

FX conversion follows **IAS 21 / ASC 830 presentation-currency translation** (design §6.5): the rate must match the *nature* of the fact, not a single as-of date.

- **Stock (instant) facts** — cash, debt, assets, liabilities, equity, PP&E — convert at the **closing spot rate** on the period-end date. These are point-in-time balances, so the point-in-time rate is correct.
- **Flow (duration) facts** — revenue, net income, operating income, D&A, capex, operating cash flow — convert at the **average rate over `[period_start, period_end]`**. These accrue across the period, so a single-day rate would misstate them (materially so for volatile currencies).

The split reuses the existing `Metric::is_instant()` classifier — no new metadata. `is_instant() == true` → closing spot; `false` → period-average.

New source — `src-tauri/src/sources/fx.rs`:

```
#[async_trait]
pub trait FxRateSource: Send + Sync {
    /// USD per 1 unit of `currency`, as micro-units (rate × 1e6), on `date`
    /// (closing spot for that day).
    async fn spot_to_usd(&self, currency: &str, date: NaiveDate)
        -> Result<Option<i64>, SourceError>;

    /// USD per 1 unit of `currency`, averaged over the inclusive date range
    /// (period-average for flow facts). Implemented as the mean of the daily
    /// spot rates the source holds in `[start, end]`.
    async fn average_to_usd(&self, currency: &str, start: NaiveDate, end: NaiveDate)
        -> Result<Option<i64>, SourceError>;
}
```

Storage — no schema change. Both methods read the *same* daily-rate cache: the existing **`fx_rate`** table keyed `(currency, date)`. `spot_to_usd` is a single-row lookup; `average_to_usd` selects the daily rows in range and means their `rate_micro`. The period-average is *derived* from daily rates, never stored as a

separate row — so there is no `fx_rate` migration. Mirror the `MarketDataSource` adapter shape and its `require_usd`-style honesty (never fabricate a rate; a gap in the daily series yields `None`, not an interpolation).

Conversion in normalization — select the rate by fact nature, then record provenance in the columns that already exist on `normalized_fact`:

```
let (rate_micro, rate_date) = if metric.is_instant() {
    (fx.spot_to_usd(ccy, period_end).await?,    period_end)
} else {
    (fx.average_to_usd(ccy, period_start, period_end).await?, period_end)
};

NormalizedFact {
    value:        usd_micro,            // home_micro × rate ÷ 1e6
    unit:         "USD".into(),
    original_value: Some(home_micro),
    original_unit:  Some(ccy),          // e.g. "INR"
    fx_rate_micro:  Some(rate_micro),   // spot or period-average, per above
    fx_rate_source: Some(source_tag),
    fx_rate_date:   Some(rate_date),    // period_end anchors both cases
    ..
}
```

`fx_rate_date` records the period-end for both cases (it anchors the conversion); the spot-vs-average distinction is implied by `Metric::is_instant()` and reproducible from the daily cache. Only if the rate lookup returns `None` do we fall through to the WP3a skip+diagnostic.

Consistency rule (design §6.5): never mix a company-reported USD convenience value with a market-rate conversion inside one derived metric. WP3a's USD-sibling preference runs first and is all-or-nothing per (metric, period); WP3b converts only the facts WP3a left in home currency, so a single derived series is either wholly convenience-USD or wholly market-converted — never a blend.

4.3 Read/UI

No IPC or UI contract change: normalized facts remain USD micro-units. The diagnostics tab already surfaces ingestion events, so WP3a skips appear there automatically.

---

5. WP-G — guards (ship with WP1)

5.1 Coverage-matrix test — `concept_map.rs` `#[cfg(test)]`

```
/// Every metric that has us-gaap candidates must also declare ifrs-full
/// candidates, unless explicitly allow-listed as US-only (bank concepts).
#[test]
fn ifrs_coverage_matches_us_gaap() {
    const US_ONLY: &[Metric] = &[
        Metric::NetInterestIncome, Metric::NoninterestIncome,
        Metric::InterestIncomeOperating, Metric::InterestExpense,
    ];
    for m in Metric::ALL {
        let cands = concepts_for(*m);
        let has_us   = cands.iter().any(|(t, _)| *t == "us-gaap");
        let has_ifrs = cands.iter().any(|(t, _)| *t == "ifrs-full");
        if has_us && !US_ONLY.contains(m) {
            assert!(has_ifrs, "metric {:?} has us-gaap but no ifrs-full candidate", m);
        }
    }
}
```

5.2 Raw-to-normalized ratio diagnostic — `orchestrator.rs`

After normalization completes, before returning, compute coverage per taxonomy and emit a WARN when a taxonomy contributes many raw facts but few normalized ones:

```
// Pseudocode — thresholds tunable; the 9/4361 INFY case must trip this.
for (taxonomy, raw_n) in raw_counts_by_taxonomy {
    let norm_n = normalized_from_taxonomy(taxonomy);
    if raw_n >= 200 && (norm_n as f64) < 0.02 * raw_n as f64 {
        record_event(deps, &cik, None, "normalize", Severity::Warn, false,
            format!("{raw_n} raw facts in `{taxonomy}` but only {norm_n} normalized — \
                     likely a concept-map coverage gap.", ),
            &mut events).await?;
    }
}
```

This converts the silent blank screen into a visible diagnostic on the very ingestion that would otherwise fail quietly.

5.3 Currency provenance assertion — normalize integration test

```
// No normalized fact may carry a non-USD unit. Any FX-converted fact must
// carry full provenance.
```

```
assert!(nf.unit == "USD");
if nf.original_unit.is_some() {
    assert!(nf.fx_rate_micro.is_some() && nf.fx_rate_date.is_some());
}
```

---

6. File-level change table

File	WP	Change
`src-tauri/src/normalize/concept_map.rs`	WP1	Append `ifrs-full` candidates to arms per §2.1; add coverage-matrix test
`src-tauri/src/normalize/concept_map.rs`	WP2	Append borrowings candidates to `LongTermDebt` / `CurrentDebt` (post sign-off)
`src-tauri/src/sources/companyfacts.rs`	WP3a	`classify_unit` / `currency_of`; keep `scale_to_micro` scaling-only
`src-tauri/src/pipeline/orchestrator.rs`	WP3a	USD-preference selection; skip+WARN for non-USD; stop hardcoding `unit:"USD"` in quarterly branch
`src-tauri/src/pipeline/orchestrator.rs`	WP-G	Raw-to-normalized ratio diagnostic (§5.2)
`src-tauri/src/sources/fx.rs`	WP3b	**New** — `FxRateSource` trait (`spot_to_usd` + `average_to_usd`) + adapter
`src-tauri/src/sources/mod.rs`	WP3b	`pub mod fx;`
`src-tauri/src/repos/fx_rate.rs`	WP3b	Read/write cache over the existing `fx_rate` table: single-row lookup for spot, range-select + mean for period-average (no schema change)
`src-tauri/src/pipeline/orchestrator.rs`	WP3b	Wire FX conversion + provenance into normalize; `IngestionDeps.fx`
`src-tauri/src/ipc/state.rs`	WP3b	Construct/inject the FX source
`src-tauri/tests/*`	WP3b	Construct `fx` in `IngestionDeps`; provenance assertion (§5.3)

No TypeScript changes for any WP (wire contract is unchanged — still USD micro-units).

---

7. Testing

7.1 WP1

- `ifrs_coverage_matches_us_gaap` test green.
- Offline accuracy test against the production DB fixture: re-normalize CIK `0001067491`; assert non-empty series for `revenue`, `net_income`, `operating_income`, `total_assets`, `total_equity`, `cash_from_operations`, `free_cash_flow`. Spot-check magnitudes (e.g. `net_income` FY2025 ≈ $3.162B, `revenue` FY2018 ≈ $10.94B).
- `cargo test` green.

7.2 WP2

- After mapping, INFY `long_term_debt` + `current_debt` ≈ 0 (no material borrowings); assert `LeaseLiabilities` is **not** picked up by either debt metric.

7.3 WP3

- WP3a: a synthetic fixture with an INR-only monetary fact and no USD sibling produces a WARN event and no normalized fact for that metric.
- WP3a: a fact with both USD and INR siblings normalizes from the USD one.
- WP3b: the same INR-only fact normalizes to USD with populated `original_*` + `fx_rate_*`; provenance assertion (§5.3) green.
- WP3b — **rate selection**: seed the daily `fx_rate` cache with a rising INR series across a fiscal year (period-average ≠ period-end spot). Assert a flow fact (e.g. `revenue`) converts at the **period-average** and a stock fact (e.g. `total_assets`) at the **closing spot** — i.e. the two use *different* `fx_rate_micro` values, matching `Metric::is_instant()`.
- WP3b — **honest gap**: with a hole in the daily cache spanning the period, `average_to_usd` / `spot_to_usd` return `None` and the fact falls through to the WP3a skip+WARN (no interpolation, no fabricated rate).

7.4 Guards

- Ratio diagnostic: re-run the pre-WP1 concept map (or a metric-stripped fixture) over INFY and assert the WARN fires; confirm it does **not** fire for a well-covered US filer.

---

8. Sequencing



All four packages ship **this cycle** (WP3b moved in 2026-07-11, design §9). The internal ordering still holds: WP1 + WP-G is the minimum to close the reported bug safely; WP2 is a one-line-per-arm change (lease policy signed off — borrowings only); WP3a establishes the USD-preference ladder that WP3b builds on. WP3b then converts the residual home-currency facts, unlocking issuers with no USD convenience translation. WP3a is a strict prerequisite for WP3b — the sibling-preference ladder decides *which* facts WP3b converts, and enforces the §4.2 consistency rule.

Implementation status: not started — deferred at the product owner's request while another agent edits source files, to avoid write conflicts.